

```
In [1]: # Step 1: Import Required Libraries
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import Adam, SGD
import cv2
```

```
In [5]: # Step 2: Load MNIST Data

# Load the dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# Shape of dataset
print("Training data shape:", X_train.shape)
print("Testing data shape:", X_test.shape)

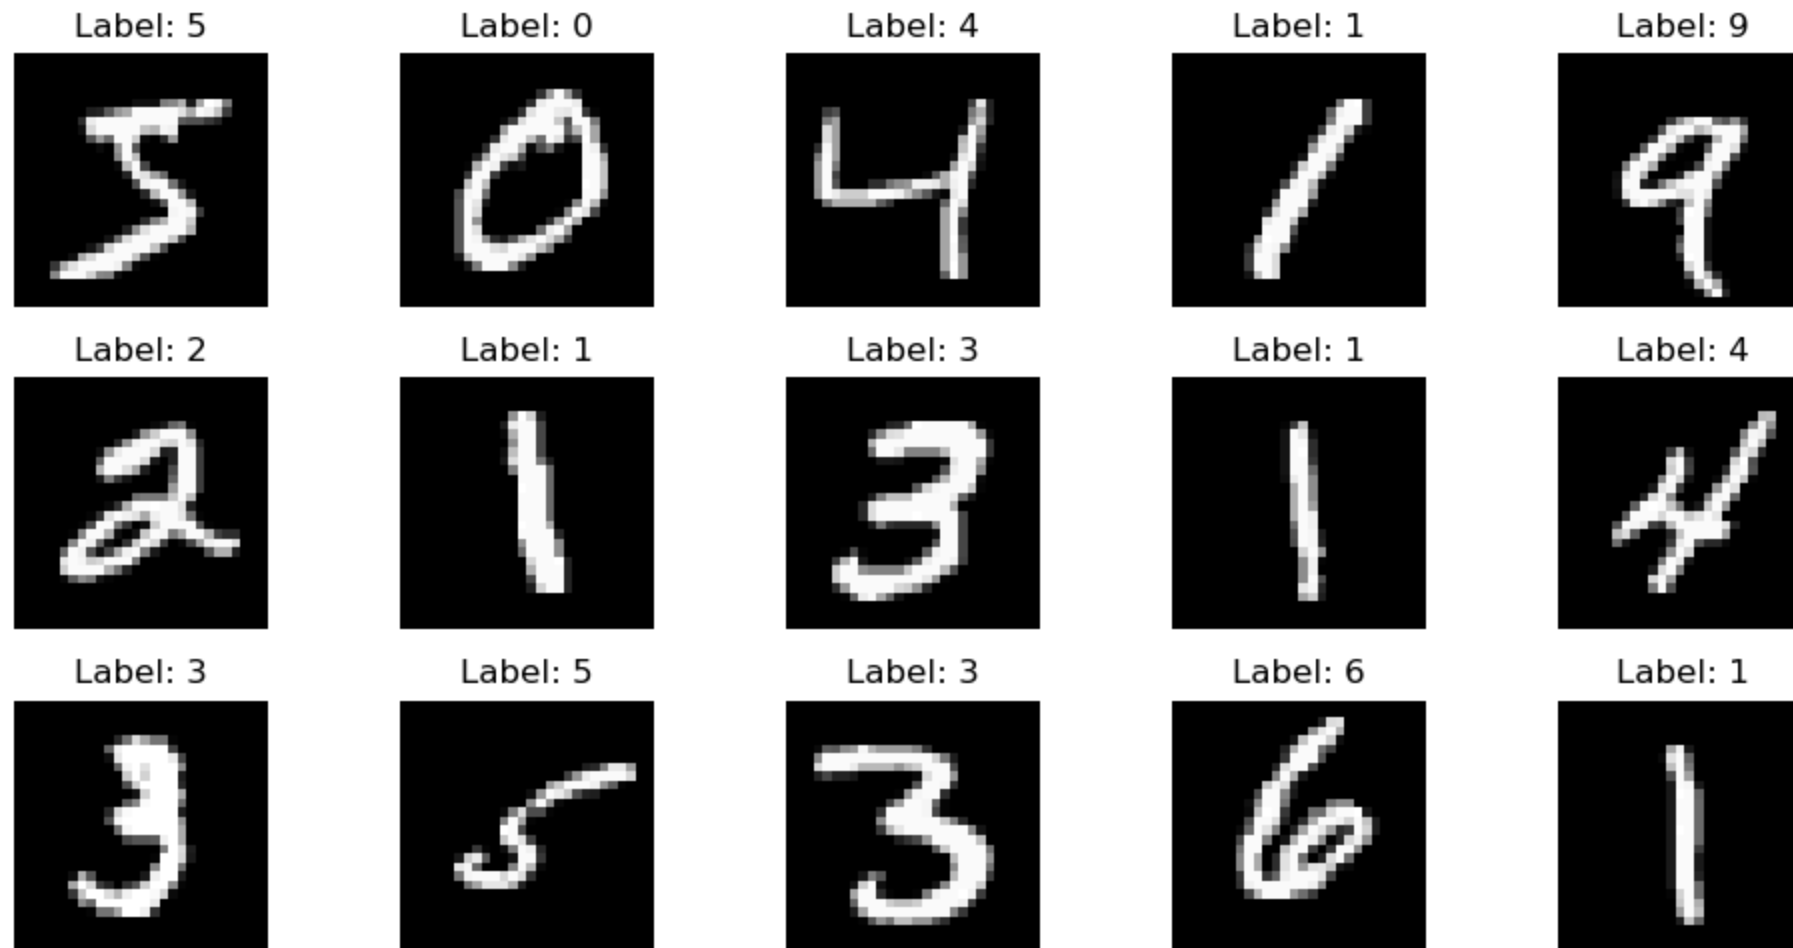
# Number of classes
print("Number of unique labels:", len(np.unique(y_train)))

# Visualize some sample digits
plt.figure(figsize=(10,5))
for i in range(15):
    plt.subplot(3,5,i+1)
    plt.imshow(X_train[i], cmap='gray')
    plt.title(f"Label: {y_train[i]}")
    plt.axis('off')
plt.tight_layout()
plt.show()
```

Training data shape: (60000, 28, 28)

Testing data shape: (10000, 28, 28)

Number of unique labels: 10



In [6]: *# Step 3A: Preprocess the Data*

```
# Normalize pixel values (0-1 range)
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# Reshape to add channel dimension (needed for CNN: 28x28x1)
X_train = X_train.reshape(-1, 28, 28, 1)
X_test = X_test.reshape(-1, 28, 28, 1)

# One-hot encode the labels
y_train = to_categorical(y_train, 10)
```

```

y_test = to_categorical(y_test, 10)

print("After preprocessing:")
print("Training data shape:", X_train.shape)
print("Testing data shape:", X_test.shape)

```

After preprocessing:

Training data shape: (60000, 28, 28, 1)

Testing data shape: (10000, 28, 28, 1)

In [7]: *# Step 3B: Build CNN Model (Function for neatness)*

```

def build_cnn_model(optimizer='adam', learning_rate=0.001):
    if optimizer == 'adam':
        opt = Adam(learning_rate=learning_rate)
    elif optimizer == 'sgd':
        opt = SGD(learning_rate=learning_rate)

    model = Sequential([
        Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),
        MaxPooling2D((2,2)),
        Conv2D(64, (3,3), activation='relu'),
        MaxPooling2D((2,2)),
        Flatten(),
        Dense(128, activation='relu'),
        Dense(10, activation='softmax')
    ])

    model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
    return model

```

In [8]: *## Train CNN Model + Plot Accuracy/Loss Graphs*

In [9]: *# Step 4A: Train the CNN Model*

```

# Build model with Adam optimizer
model_adam = build_cnn_model(optimizer='adam', learning_rate=0.001)

# Train the model
history_adam = model_adam.fit(
    X_train, y_train,
    validation_split=0.2,

```

```

    epochs=15,
    batch_size=100,
    verbose=2
)

```

Epoch 1/15

/opt/anaconda3/lib/python3.12/site-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
480/480 - 6s - 12ms/step - accuracy: 0.9344 - loss: 0.2238 - val_accuracy: 0.9768 - val_loss: 0.0793
Epoch 2/15
480/480 - 6s - 12ms/step - accuracy: 0.9811 - loss: 0.0603 - val_accuracy: 0.9865 - val_loss: 0.0490
Epoch 3/15
480/480 - 6s - 13ms/step - accuracy: 0.9871 - loss: 0.0395 - val_accuracy: 0.9851 - val_loss: 0.0507
Epoch 4/15
480/480 - 6s - 12ms/step - accuracy: 0.9908 - loss: 0.0290 - val_accuracy: 0.9902 - val_loss: 0.0365
Epoch 5/15
480/480 - 6s - 12ms/step - accuracy: 0.9928 - loss: 0.0229 - val_accuracy: 0.9891 - val_loss: 0.0361
Epoch 6/15
480/480 - 6s - 13ms/step - accuracy: 0.9938 - loss: 0.0192 - val_accuracy: 0.9883 - val_loss: 0.0388
Epoch 7/15
480/480 - 6s - 13ms/step - accuracy: 0.9957 - loss: 0.0137 - val_accuracy: 0.9909 - val_loss: 0.0368
Epoch 8/15
480/480 - 6s - 12ms/step - accuracy: 0.9964 - loss: 0.0108 - val_accuracy: 0.9904 - val_loss: 0.0370
Epoch 9/15
480/480 - 6s - 13ms/step - accuracy: 0.9967 - loss: 0.0097 - val_accuracy: 0.9902 - val_loss: 0.0413
Epoch 10/15
480/480 - 6s - 13ms/step - accuracy: 0.9969 - loss: 0.0083 - val_accuracy: 0.9891 - val_loss: 0.0449
Epoch 11/15
480/480 - 7s - 14ms/step - accuracy: 0.9974 - loss: 0.0076 - val_accuracy: 0.9896 - val_loss: 0.0410
Epoch 12/15
480/480 - 6s - 13ms/step - accuracy: 0.9982 - loss: 0.0054 - val_accuracy: 0.9903 - val_loss: 0.0468
Epoch 13/15
480/480 - 7s - 14ms/step - accuracy: 0.9981 - loss: 0.0055 - val_accuracy: 0.9897 - val_loss: 0.0474
Epoch 14/15
480/480 - 7s - 14ms/step - accuracy: 0.9990 - loss: 0.0031 - val_accuracy: 0.9897 - val_loss: 0.0483
Epoch 15/15
480/480 - 6s - 13ms/step - accuracy: 0.9980 - loss: 0.0060 - val_accuracy: 0.9912 - val_loss: 0.0457

```

In [10]: *# Step 4B: Plot Training vs Validation Accuracy and Loss*

```
def plot_training_curves(history, title='Model'):
```



```
plt.figure(figsize=(12,5))

# Accuracy Plot
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title(f'{title} - Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Loss Plot
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title(f'{title} - Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

# Now plot for Adam optimizer model
plot_training_curves(history_adam, title='CNN with Adam Optimizer')
```



```
In [11]: # Step 5: Evaluate the trained CNN Model

from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns

def evaluate_model(model, X_test, y_test):
    # Predict probabilities
    y_pred_probs = model.predict(X_test)

    # Convert probabilities to class labels
    y_pred = np.argmax(y_pred_probs, axis=1)
    y_true = np.argmax(y_test, axis=1)
```


```
# Confusion Matrix
cm = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

# Classification Report
print("\nClassification Report:")
print(classification_report(y_true, y_pred))

# Now evaluate the model you trained (Adam optimizer model)
evaluate_model(model_adam, X_test, y_test)
```

313/313  1s 2ms/step





Classification Report:				
	precision	recall	f1-score	support
0	0.99	1.00	0.99	980
1	1.00	1.00	1.00	1135
2	0.99	1.00	0.99	1032
3	0.99	1.00	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.99	0.99	892
6	1.00	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.99	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

In [12]: *# Step 6A: Build and Train another CNN Model using SGD Optimizer*

```
# Build model with SGD optimizer
model_sgd = build_cnn_model(optimizer='sgd', learning_rate=0.01)

# Train the model
history_sgd = model_sgd.fit(
    X_train, y_train,
    validation_split=0.2,
    epochs=15,
    batch_size=100,
    verbose=2
)
```



Epoch 1/15

/opt/anaconda3/lib/python3.12/site-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

480/480 - 6s - 13ms/step - accuracy: 0.7183 - loss: 1.0462 - val\_accuracy: 0.8974 - val\_loss: 0.3477  
Epoch 2/15  
480/480 - 6s - 13ms/step - accuracy: 0.9072 - loss: 0.3123 - val\_accuracy: 0.9266 - val\_loss: 0.2445  
Epoch 3/15  
480/480 - 6s - 13ms/step - accuracy: 0.9325 - loss: 0.2283 - val\_accuracy: 0.9481 - val\_loss: 0.1829  
Epoch 4/15  
480/480 - 6s - 13ms/step - accuracy: 0.9461 - loss: 0.1813 - val\_accuracy: 0.9584 - val\_loss: 0.1520  
Epoch 5/15  
480/480 - 6s - 13ms/step - accuracy: 0.9555 - loss: 0.1500 - val\_accuracy: 0.9638 - val\_loss: 0.1307  
Epoch 6/15  
480/480 - 6s - 13ms/step - accuracy: 0.9615 - loss: 0.1284 - val\_accuracy: 0.9658 - val\_loss: 0.1206  
Epoch 7/15  
480/480 - 6s - 13ms/step - accuracy: 0.9661 - loss: 0.1135 - val\_accuracy: 0.9688 - val\_loss: 0.1098  
Epoch 8/15  
480/480 - 6s - 13ms/step - accuracy: 0.9695 - loss: 0.1015 - val\_accuracy: 0.9724 - val\_loss: 0.0973  
Epoch 9/15  
480/480 - 6s - 13ms/step - accuracy: 0.9719 - loss: 0.0937 - val\_accuracy: 0.9723 - val\_loss: 0.0955  
Epoch 10/15  
480/480 - 6s - 13ms/step - accuracy: 0.9747 - loss: 0.0859 - val\_accuracy: 0.9752 - val\_loss: 0.0843  
Epoch 11/15  
480/480 - 6s - 13ms/step - accuracy: 0.9755 - loss: 0.0802 - val\_accuracy: 0.9734 - val\_loss: 0.0877  
Epoch 12/15  
480/480 - 6s - 13ms/step - accuracy: 0.9772 - loss: 0.0741 - val\_accuracy: 0.9755 - val\_loss: 0.0827  
Epoch 13/15  
480/480 - 6s - 13ms/step - accuracy: 0.9788 - loss: 0.0707 - val\_accuracy: 0.9775 - val\_loss: 0.0755  
Epoch 14/15  
480/480 - 6s - 13ms/step - accuracy: 0.9798 - loss: 0.0667 - val\_accuracy: 0.9787 - val\_loss: 0.0757  
Epoch 15/15  
480/480 - 6s - 13ms/step - accuracy: 0.9809 - loss: 0.0627 - val\_accuracy: 0.9777 - val\_loss: 0.0744


In [13]: *# Step 6B: Plot Training vs Validation Accuracy and Loss for SGD Model*

```
plot_training_curves(history_sgd, title='CNN with SGD Optimizer')
```



```
In [14]: # Step 6C: Evaluate the SGD Optimizer Model
         evaluate_model(model_sgd, X_test, y_test)
```

313/313 ————— 1s 2ms/step



Classification Report:				
	precision	recall	f1-score	support
0	0.97	0.99	0.98	980
1	0.99	0.99	0.99	1135
2	0.99	0.98	0.98	1032
3	0.98	0.98	0.98	1010
4	0.98	0.99	0.99	982
5	0.96	0.99	0.97	892
6	0.99	0.98	0.98	958
7	0.98	0.98	0.98	1028
8	0.99	0.94	0.96	974
9	0.97	0.98	0.98	1009
accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

In [15]: *# Step 7: Predict a Single Digit Image using the trained model*

```
import cv2

def predict_single_digit(image_path, model):
    # Load image in grayscale
    img = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)

    # Resize to 28x28
    img = cv2.resize(img, (28, 28))

    # Invert colors if needed
    img = 255 - img

    # Normalize pixel values
    img = img.astype('float32') / 255.0

    # Reshape to model input shape (1,28,28,1)
    img = img.reshape(1, 28, 28, 1)

    # Predict
    prediction = model.predict(img)
    predicted_digit = np.argmax(prediction)
```

```
# Show image and prediction
plt.imshow(img.squeeze(), cmap='gray')
plt.title(f"Predicted Digit: {predicted_digit}")
plt.axis('off')
plt.show()


return predicted_digit
```

```
In [16]: # Give the path of your custom digit image
image_path = '/Users/manoharshasappa/Desktop/Files/3 Sem/AIT 736 NLP/Final_Project/Screenshot 2025-04-26 at 11.13.57 AM.png'

# Predict using Adam model
predicted_digit = predict_single_digit(image_path, model_adam)
print(f"Predicted Digit: {predicted_digit}")
```

1/1 ————— 0s 13ms/step

Predicted Digit: 8



Predicted Digit: 8

In [ ]:


In this project, we built a CNN model to classify handwritten digits using the MNIST dataset. We trained two models using Adam and SGD optimizers for comparison. Data preprocessing steps included normalization and reshaping. The models were evaluated using accuracy/loss curves, confusion matrices, and classification reports. Hyperparameter tuning was performed by adjusting learning rate and batch size. The final system successfully predicts single handwritten digits from custom images with high accuracy

```
In [18]: image_path = '/Users/manoharshasappa/Desktop/Files/3 Sem/AIT 736 NLP/Final_Project/Screenshot 2025-04-26 at 12.08.46 PM.png'

# Predict using Adam model
predicted_digit = predict_single_digit(image_path, model_adam)
print(f"Predicted Digit: {predicted_digit}")
```

1/1 ————— 0s 11ms/step

Predicted Digit: 4




Predicted Digit: 4

```
In [26]: image_path = '/Users/manoharshasappa/Desktop/Files/3 Sem/AIT 736 NLP/Final_Project/Screenshot 2025-04-27 at 11.10.48 PM.png'
```

```
# Predict using Adam model
predicted_digit = predict_single_digit(image_path, model_adam)
print(f"Predicted Digit: {predicted_digit}")
```

1/1 ————— 0s 12ms/step

Predicted Digit: 9



Predicted Digit: 9

```
In [ ]:
```